PFA: Portable Format for Analytics

Jim Pivarski

Sometime in 2014

Abstract

This specification defines the syntax and semantics of the Portable Format for Analytics (PFA).

PFA is a mini-language for mathematical calculations that is usually generated programmatically, rather than by hand. A PFA document is a string of JSON-formatted text that describes an executable called a scoring engine. Each engine has a well-defined input, a well-defined output, and functions for combining inputs to construct the output in an expression-centric syntax tree. In addition, it has centralized facilities for maintaining state, with well-defined semantics for sharing state among scoring engines in a thread-safe way. The specification defines a suite of mathematical and statistical functions for transforming data, but it does not define any means of communication with an operating system, file system, or network. A PFA engine must be embedded in a larger system that has these capabilities, and thus an analytic workflow is decoupled into a part that manages data pipelines (such as Hadoop, Storm, or Akka), and a part that describes the algorithm to be performed on data (PFA).

PFA is similar to the Predictive Model Markup Language (PMML), an XML-based specification for statistical models, but whereas PMML's focus is on statistical models in the abstract, PFA's focus is on the scoring procedure itself. The same input given to two PFA-enabled systems must yield the same output, regardless of platform (e.g. a JVM in Hadoop, a client's web browser, a GPU kernel function, or even an IP core directly embedded in an integrated circuit). Unlike PMML, the PFA specification defines the exact bit-for-bit behavior of any well-formed document, the semantics of data types and data structures, including behavior in concurrent systems, and all cases in which an exception should be thrown. Like PMML, PFA is a specification, not an implementation, it defines a suite of statistical algorithms for analyzing data, and it is usually generated programmatically, as the output of a machine learning algorithm, for instance.

Status of this document

This section describes the status of this document at the time of the current draft. Other documents may supersede this document.

This document is an early draft that has not been endorsed for recommendation by any organization. It describes a proposed specification that could, in the future, become a standard.

Contents

1 Introduction

1.1 Motivation for PFA

The Portable Format for Analytics (PFA) is a mini-language for mathematical calculations. It differs from most programming languages in that it is optimized for automatic code generation, rather than writing programs by hand. The primary use-case is to represent the output of machine learning algorithms, such that they can be freely moved between systems. Traditionally, this field has been dominated by special-purpose file formats, each representing only one type of statistical model. The Predictive Model Markup Language (PMML) provides a means of unifing the most common model types into one file format. However, PMML can only express a fixed set of pre-defined model types; new model types must be agreed upon by the Data Mining Group (DMG) and integrated into a new version of PMML, then that new version must be adopted by the community before it is widely usable.

PFA represents models and analytic procedures more generally by providing generic programming constructs, such as conditionals, loops, persistent state, and callback functions, in addition to a basic suite of statistical tools. Conventional models like regression, decision trees, and clustering are expressed by referencing the appropriate library function, just as in PMML, but new models can be expressed by composing library functions or passing user-defined callbacks. Most new statistical techniques are variants of old techniques, so a small number of functions with the appropriate hooks for inserting user code can represent a wide variety of methods, many of which have not been discovered yet.

Given that flexibility is important, one might consider using a general purpose programming language, such as C, Java, Python, or especially R, which is specifically designed for statistics. While this is often the easiest method for small problems that are explored, formulated, and solved on an analyst's computer, it is difficult to scale up to network-sized solutions or to deploy on production systems that need to be more carefully controlled than a personal laptop. The special-purpose code may depend on libraries that cannot be deployed, or may even be hard to identify exhaustively. In some cases, the custom code might be regarded as a stability or security threat that must be thoroughly reviewed before deployment. If the analytic algorithm needs to be deployed multiple times before it is satisfactory and each deployment is reviewed for reasons unrelated to its analytic content, development would be delayed unnecessarily. This problem is solved by decoupling the analytic workflow into a part that deals exclusively with mathematics (the PFA scoring engine) and the rest of the infrastructure (the PFA host). A mathematical algorithm implemented in PFA can be updated frequently with minimal review, since PFA is incapable of raising most stability or security issues, due to its limited access.

PFA is restricted to the following operations: mathematical functions on numbers, strings, raw bytes, homogeneous lists, homogeneous maps (also known as hash-tables, associative arrays, or dictionaries), heterogeneous records, and unions of the above, where mathematical functions include basic operations, special functions, data sturcture manipulations, missing data handling, descriptive statistics, and common model types such as regression, decision trees, and clustering, parameterized for flexibility. PFA does not include any means of accessing the operating system, the file system, or the network, so a rouge PFA engine cannot expose or manipulate data other than that which is intentionally funneled into it by the host system. The full PFA specification allows recursion and unterminated loops, but execution time is limited by a timeout. PFA documents may need to be reviewed for mathematical correctness, but they do not need to be reviewed for safety.

Another reason to use PFA as an intermediate model representation is for simplicity of code generation. A machine learning algorithm generates an executable procedure, usually a simple, parameterized decider algorithm that categorizes or makes predictions based on new data. Although the parameters might be encoded in a static file, some component must be executable. A PFA document bundles the executable with its parameters, simplifying version control.

The syntax of PFA is better suited to automatic code generation than most programming languages.

Many languages have complex syntax to accomodate the way people think while programming, including infix operators, a distinction between statements and expressions, and in some cases even meaningful whitespace. Though useful when writing programs by hand, these features only complicate automatic code generation. A PFA document is an expression tree rendered in JSON, and trees are easy to programmatically compose into larger trees without introducing syntax errors in the generated code. This is well-known in the Lisp community, since the ease of writing code-modifying macros in Lisp is often credited to its exclusive use of expression trees, rendered as parenthesized lists (known as S-expressions). PFA uses JSON, rather than S-expressions, because libraries for manipulating JSON objects are more widely available and JSON provides a convenient syntax for maps, but the transliteration between JSON and S-expressions is straight-forward.

Another benefit of PFA's simplicity relative to general programming languages is that it is more amenable to static analysis. A PFA host can more thoroughly examine an incoming PFA document for undesirable features. Although PFA makes use of callback functions to provide generic algorithms, functions are not first-class objects in the language, meaning that they cannot be dynamically assigned to variables. The identity of every function call can be determined without running the engine, which makes it possible to statically generate a graph of function calls and identify recursive loops. In very limited runtime environments, such as some GPUs, the compiler implicitly inlines all function calls, so recursion is not possible. In cases like these, static analysis of the PFA document is a necessary step in generating the executable.

A PFA document can also be statically type-checked. This allows for faster execution times, since types do not need to be checked at run-time, but it also provides additional safety to the PFA host.

PFA uses Apache Avro schemae for type annotations. Avro is an open-source serialization protocol, widely used in Hadoop and related projects, whose type schemae are expressed as JSON objects and whose data structures can be expressed as JSON objects. Therefore, all parts of the PFA engine, including control structures, type annotations, and embedded data are all expressed in one seamless JSON object. Avro additionally has well-defined rules to resolve different but possibly compatible schemae, which PFA reinterprets as type promotion (allowing integers to be passed to a function that expects floating-point numbers, for instance). When interpreted this way, Avro also has a type-safe null, which PFA uses to ensure that missing data are always explicitly handled. Finally, the input and output of every PFA engine can always be readily (de)serialized into Avro's binary format or JSON representation, since Avro libraries are available on a wide variety of platforms.

1.2 Terminology used in this specification

Within this specification, the key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "MAY", and "OPTIONAL" are to be interpreted as described in RFC 2119 (see RFC2119). However, for readability, these words do not appear in all uppercase letters in this specification.

At times, this specification provides hints and suggestions for implementation. These suggestions are not normative and conformance with this specification does not depend on their realization. These hints contain the expression "We suggest…", "Specific implementations may…", or similar wording.

This specification uses the terms "JSON object", "JSON object member name", "JSON object member value", "JSON array", "JSON array value", "number", "integer", "string", "boolean", and "null" as defined in the JSON specification (RFC-4627), sections 2.2 through 2.5. It also references and quotes sections of the Avro 1.7.6 specification (`http://avro.apache.org/docs/1.7.6/spec.html`).

1.3 PFA MIME type and file name extension

The recommended MIME type for PFA is "application/pfa+json", though this is not yet in the process of standardization.

It is recommended that PFA files have the extension ".pfa" (all lowercase) on all platforms. It is recommended that gzip-compressed PFA files have the extension ".pfaz" (all lowercase) on all platforms.

1.4 Levels of PFA conformance and PFA subsets

PFA is a large specification with many modules, so some projects or vendors may wish to implement some but not all of the specification. However, interoperability is the reason PFA exists; if an implementation does not adhere to the standard, it has limited value. It is therefore useful to explicitly define what it means for a system to partially implement the standard.

JSON subtrees of a PFA document are interpreted in the following six contexts.

- Top-level fields are JSON object member name, value pairs in the outermost JSON object of the PFA document. They have unique member names and describe global aspects of the scoring engine.
- Special forms are JSON objects that specify executable expressions and function definitions. Each is associated with a unique name.
- Library functions are strings that specify routines not defined in the PFA document itself. Each is associated with a unique name that does not conflict with any of the special forms' names.
- Avro type schemae are JSON objects and strings that describe data types. The syntax and meaning of Avro types are specified in the Avro 1.7.6 specification.
- Embedded data are JSON objects, JSON arrays, numbers, integers, strings, booleans, and nulls that describe data structures. The syntax and meaning of these objects are also defined by Avro, as the format used by the `JSONEncoder` and `JSONDecoder`.
- Options are JSON object member values of the `options` top-level field and may be overridden by the PFA system. They all have well-defined defaults and unique, hierarchical names.

A system may be partially PFA compliant if it implements some but not all top-level fields, some but not all special forms, or some but not all library functions. Its coverage may be specified by listing the object member names of the top-level fields that it does implement, the names of the special forms that it does implement, and the names of the library functions that it does implement. Those top-level fields, special forms, and library functions that it does implement must be completely and correctly implemented. The coverage is therefore atomic and one can immediately determine if a particular system can execute a particular PFA document by checking the set of names used by the document against the set of names implemented by the system.

Some special forms and library functions make use of some top-level fields. For example, library functions that generate random numbers use the `randseed` field for configuration. These special forms and library functions cannot be considered implemented unless the corresponding top-level fields are also implemented. The dependencies are explicitly defined in this specification.

Avro type schemae and JSON-encoded data should be completely implemented, to the extent defined by the Avro specification. We suggest that implementations use language-specific Avro libraries as much as possible, rather than implementing Avro-related features in a PFA system.

Options may also be implemented atomically by name. If a named option is not implemented, the system should behave as though that option had its default value, regardless of whether the option is explicitly set in the PFA document. Options can in general be overridden by a host system, so if a host system doesn't implement an option, it is as though the system enforces the default.

The PFA standard is defined so that a PFA-compliant system can verify that the JSON types of a PFA document are correctly composed (syntax check), verify that the PFA invariants are maintained and Avro

data types are correctly composed (semantic check), and impose additional constraints on the set of top-level fields, special forms, and library functions used (optional checks). A PFA-compliant system should perform the syntax and semantic checks, including all type inference and type checking, though it is not strictly required. A PFA document that does not satisfy these invariants and type constraints is not valid and its behavior is not defined by this specification. The third set of checks, however, is completely optional and different systems may apply different constraints on the kinds of scoring engines they are willing to execute. For instance, an implementation targeting a limited environment in which recursion is not possible may analyze the document and reject it if any recursive loops are found.

This specification does not define any standardized subsets of PFA. As stated above, partial conformance is defined by ad hoc subsets of atomic units. However, as experience develops, the community may define industry-standard subsets of PFA for specific purposes or special environments. Conforming to a standardized subset would provide better interoperability than defining ad hoc subsets, and we would recommend such a standard when it exists. At present, we can only recommend a carefully chosen ad hoc subset or complete conformance.

2 PFA document structure

A PFA document is a serialized JSON object representing an executable scoring engine. Only the following JSON object member names may appear at this JSON nesting level. These are the top-level fields referred to in the conformance section of this specification. Three fields, **`action`**, **`input`**, and **`output`**, are required for every PFA document and are therefore required for every PFA implementation. The rest are optional for PFA documents and not strictly required for PFA implementations. As explained in the conformance section, not implementing some top-level fields can make some special forms and functions unimplementable.

`name`: A string used to identify the scoring engine (has no effect on calculations).

`method`: A string that may be "map", "emit", or "fold" (see Sec. 3.2). If absent, the default value is "map".

`input`: An Avro schema representing the data type of data provided to the scoring engine (see Sec. 3.3).

`output`: An Avro schema representing the data type of data produced by the scoring engine (see Sec. 3.3). The way that output is returned to the host system depends on the **`method`**.

`begin`: An expression or JSON array of expressions that are executed in the begin phase of the scoring engine's run (see Sec. 3.1).

`action`: An expression or JSON array of expressions that are executed for each input datum in the active phase of the scoring engine's run (see Sec. 3.1).

`end`: An expression or JSON array of expressions that are executed in the end phase of the scoring engine's run (see Sec. 3.1).

`fcns`: A JSON object whose member values are function definitions, defining routines that may be called by expressions in **`begin`**, **`action`**, **`end`**, or by expressions in other functions.

`zero`: Embedded JSON data whose type must match the **`output`** type of the engine. This is only used by the "fold" method to initialize the fold aggregation. If **`method`** is "map" or "emit", this field is ignored.

`cells`: A JSON object whose member values specify statically allocated, named, typed units of persistent state or embedded data (see Sec. 3.4). The format of this JSON object is restricted: see Sec. 2.1.

`pools`: A JSON object whose member values specify dynamically allocated namespaces of typed persistent state (see Sec. 3.4). The format of this JSON object is restricted: see Sec. 2.1.

`randseed`: An integer which, if present, sets the seed for pseudorandom number generation (see Sec. 3.8).

`doc`: A string used to describe the scoring engine or its provenance (has no effect on calculations).

`metadata`: A JSON object, array, string, number, boolean, or null used to describe the scoring engine or its provenance (has no effect on calculations).

`options`: A JSON object of JSON objects, arrays, strings, numbers, booleans, or nulls used to control execution. The set of possible options and their representation is restricted:: see see Sec. 3.7.

Example 2.1. This is the simplest possible PFA document. It only reads `null` values, returns `null` values, and performs no calculations.

```
{"input": "null", "output": "null", "action": null}
```

Example 2.2. This is a simple yet non-degenerate PFA document. It increments numerical input by 1.

```
{"input": "double", "output": "double", "action": {"+": ["input", 1]}}
```

Example 2.3. This example implements a small decision tree. Input data are records with three fields: "one" (integer), "two" (double), and "three" (string). The decision tree is stored in a cell named "tree" with type "TreeNode". The tree has three binary splits (four leaves). The scoring engine walks from the root to a leaf for each input datum, choosing a path based on values found in the record's fields, and returns the string it finds at the tree's leaf. (See the definition of the model.tree.simpleWalk function.)

```
{"input": {"type": "record", "name": "Datum", "fields":
   [{"name": "one", "type": "int"},
    {"name": "two", "type": "double"},
    {"name": "three", "type": "string"}]},
 "output": "string",
 "cells": {"tree":
             {"type":
               {"type": "record",
                "name": "TreeNode",
                "fields": [
                  {"name": "field", "type": "string"},
                  {"name": "operator", "type": "string"},
                  {"name": "value", "type": ["double", "string"]},
                  {"name": "pass", "type": ["string", "TreeNode"]},
                  {"name": "fail", "type": ["string", "TreeNode"]}]},
              "init":
                {"field": "one",
                 "operator": "<",
                 "value": {"double": 12},
                 "pass":
                   {"TreeNode":
                     {"field": "two",
                      "operator": ">",
                      "value": {"double": 3.5},
                      "pass": {"string": "yes-yes"},
                      "fail": {"string": "yes-no"}}},
                 "fail":
                   {"TreeNode":
                     {"field": "three",
                      "operator": "==",
                      "value": {"string": "TEST"},
                      "pass": {"string": "no-yes"},
                      "fail": {"string": "no-no"}}}}}},
 "action":
   {"model.tree.simpleWalk": ["input", {"cell": "tree"}]}}
```

2.1 Cells and Pools

The **cells** and **pools** top-level fields, if present, are JSON objects whose member values are cell-specifications or pool-specifications, respectively. A cell is a mutable, global data store that holds a single value with a specific type, and a pool is a mutable map from dynamically allocated names to values of a specific type (see Sec. 3.4).

A cell-specification is a JSON object with the following fields.

type: *(required)* An Avro schema representing the data type of this cell.

init: *(required)* Embedded JSON whose type must match **type**. This is the initial value of the cell (or constant value if it is never modified).

shared: An optional boolean specifying whether this cell is thread-local to one scoring engine or shared among a battery of similar engines (see Sec. 3.5). The default is **false**.

rollback: An optional boolean specifying whether this cell should be rolled back to the state it had at the beginning of an **action** if an exception occurs during the **action**. The default is **false**, and **shared** and **rollback** are mututally incompatible: they cannot both be **true**.

A pool-specification is a JSON object with the following fields.

type: *(required)* An Avro schema representing the data type of this pool.

init: JSON object whose member values are embedded JSON that must match **type**. Unlike a cell, a pool may be empty on initialization, in which case **init** is either unspecified or **{}**.

shared: An optional boolean specifying whether this pool is thread-local to one scoring engine or shared among a battery of similar engines (see Sec. 3.5). The default is **false**.

rollback: An optional boolean specifying whether this pool should be rolled back to the state it had at the beginning of an **action** if an exception occurs during the **action**. The default is **false**, and **shared** and **rollback** are mututally incompatible: they cannot both be **true**.

A complete explanation of cells and pools is given in Sec. 3.4.

3 Scoring engine execution model

A PFA document (string of JSON-formatted text) describes a PFA scoring engine (executable routine) or a battery of initially identical engines. An engine behaves as a single-threaded executable with global state (cells and pools) and local variables. A battery of scoring engines may run in parallel and only share data if some cells or pools are explicitly marked as **`shared`**. Although a battery of scoring engines generated by a single PFA document start in exactly the same state, they may evolve into different states if they have any unshared cells or pools.

PFA engines are units of work that may fit into a pipeline system like Hadoop, Storm, or Akka. In a map-reduce framework such as Hadoop, for instance, one PFA document could describe the calculation performed by all of the mappers and another could describe the calculation performed by all of the reducers. The mappers are a battery of independent PFA engines, as are the reducers. In pure map-reduce, the mappers would not communicate with each other and the reducers would not communicate with each other, so none of the cells or pools should be marked as **`shared`**. With this separation of concerns, issues of transferring data, interpreting input file types, and formatting output should be handled by the pipeline system (Hadoop in this case) while the mathematical procedure is handled by PFA. Changing file formats would require an update to the pipeline code (and possibly a code review), but changing details of the analytic would only require a new PFA document (a JSON configuration file).

3.1 Execution phases of a PFA scoring engine

A PFA engine has a 7 phase lifecycle. These phases are the following, executed in this order:

1. reading the PFA document and performing a syntax check;
2. verifying PFA invariants and checking type consistency;
3. additional checks, constraints required by a particular PFA system;
4. initialization of the engine;
5. execution of the **`begin`** routine;
6. execution of the **`action`** routine for each input datum;
7. execution of the **`end`** routine.

In phase 1, JSON is decoded and may be used to build an abstract syntax tree of the whole document. At this stage, JSON types must be correctly matched (e.g. if a number is expected, a string cannot be provided instead) to build the syntax tree. Incorrectly formatted JSON should also be rejected, though we recommend that a dedicated JSON decoder is used for this task. Avro schemae should also be interpreted in this phase (see Sec. 4.1).

In phase 2, the loaded PFA document is interpreted as an executable. If the specific PFA implementation builds code with macros, compiles bytecode, or synthesizes a circuit for execution, that work should happen in this phase. Data types should be inferred and checked (see Sec. 4.2), especially if the executable is compiled.

Phase 3 is provided for optional checks. Due to limitations of a particular environment, some PFA systems may need to be more restrictive than the general specification and reject what would otherwise be a valid PFA document. Reasons include unimplemented function calls, inability to implement recursion, or data structures that are too large. The phase 3 checks may need to be performed concurrently with the phase 2 checks to build the executable.

Phase 4, initialization, is when data structures such as cells and pools are allocated and filled, network connections are established (if relevant for a particular PFA implementation), pseudorandom number generators are seeded, etc. These are actions that the engine must perform to work properly but are not a part of the **begin**, **action**, or **end** routines.

The actions performed in the last three phases, **begin**, **action**, and **end**, are explicitly defined in the PFA document. A PFA system must implement the **action** phase, since every PFA document must define an **action**. The **action** accepts input and returns output, though the way it does so depends on the **method** (Sec. 3.2).

The **begin** and **end** phases do not accept input and do not return output: they can only modify cells and pools, emit log messages, or throw exceptions. A PFA system is not required to implement **begin** and **end**. If a system that does not implement **begin** encounters a document that has a **begin** routine, it must fail with an error. If a system that does not implement **end** encounters a document that has an **end** routine, it need not fail with an error, though it may. This is because some PFA documents may use **begin** to initialize essential data structures and the **action** would only function properly if **begin** has been executed, but the **end** routine can only affect the state of a completed scoring engine whose interpretation is implementation-specific. Moreover, some data pipelines do not even have a concept of completion, such as Storm.

After all input data have been passed to the scoring engine and the last **action** or **end** routine has finished, the scoring engine is said to be completed. This may be considered an eighth phase of the engine, though its behavior at this point is not defined by this specification. A particular PFA system may extract aggregated results from a completed engine's state and it may even call functions defined in the document's **fcn** field, but this is beyond the scope of the standard PFA lifecycle. (Note: if the primary purpose of a scoring engine is to aggregate data, consider using the "fold" **method** instead of extracting from the engine's internal state.)

A completed scoring engine may be used to create a new PFA document, in which the final state of the cells and pools are used to define the **cell init** or **pool init** of the new document, such that a new scoring engine would start where the old one left off. A PFA system may even re-use an old scoring engine as a new scoring engine (repeating phase 4 onward), but a re-used engine must behave exactly like a new engine with copied state, such that the re-use is an implementation detail and does not affect behavior.

A PFA system may call functions defined in the document's **fcn** field at any time, but if the function modifies state (a "cell-to" or "pool-to" special form is reachable in its call graph) and the engine is not complete, the function call must not be allowed because it could affect the engine's behavior. A PFA system must not execute **begin**, **action**, or **end** outside of its lifecycle.

3.2 Scoring method: map, emit, and fold

PFA defines the following three methods for calling the **action** routine of a scoring engine.

"map": The **action** routine is given an **input** value, which it uses to construct and return an output. Barring exceptions, the output dataset would have exactly as many values as the input dataset.

"emit": The **action** routine is given an **input** value and an **emit** callback function, and the functional return value is ignored. The scoring engine returns results to the host system by calling **emit**. It can call **emit** any number of times, and thus the output dataset may be smaller or larger than the input dataset. For example, a filter would call **emit** zero or one times for each input.

"fold": The **action** routine is given an **input** value and a **tally** value, which it uses to construct and return an output. The first time **action** is invoked, **tally** is equal to **zero** (the top-level field). On the N^{th} time **action** is invoked, **tally** is equal to the $(N-1)^{\text{th}}$ return value. Thus, a "fold" scoring engine is an aggregator: transformed inputs may be counted, summed, maximized, or otherwise accumulated in the **tally**. The aggregate of the entire input dataset is the last return value of **action**, though the host system may make use of partial sums as well.

For all three methods, the **input** is available to expressions as a read-only symbol that can be accessed in an expression as the JSON string **"input"** (see Sec. 7.2). The **input** symbol's scope is limited to the **action** routine: it is not accessible in user-defined functions unless explicitly passed. The **input** symbol's data type is specified by the top-level field named **input**.

For the "map" and "fold" methods, the data type of the last expression in the **action** routine must be the type specified by the top-level field named **output**. For the "emit" method, there is no constraint on the type of the last expression in **action**, but the argument passed to the **emit** function must have **output** type.

For the "emit" method, the **emit** function is a globally accessible function. It may be called or referenced without qualification by any user-defined function or even in the **begin** and **end** routines.

For the "fold" method, the **tally** is available to expressions as a read-only symbol that can be accessed in an expression as the JSON string **"tally"** (see Sec. 7.2). The **tally** symbol's scope is limited to the **action** routine: it is not accessible in user-defined functions unless explicitly passed. The **tally** symbol's data type is specified by the **output** top-level field. The top-level field named **zero** must also have **output** type.

The means by which input values are provided to the scoring engine, output values are retrieved, and the **emit** function is set or changed are all unspecified. A PFA system may change the **emit** function at any time, even while an **action** is being processed (though we do not recommend this). However, the **emit** function must be defined and callable at all times during the **begin**, **action**, and **end** phases of the scoring engine's lifecycle.

3.3 Input and output type specification

The member values of the top-level fields **input** and **output** are Avro schemae (see Sec. 4.1). The way that these types constrain the input and output of scoring engines depends on the **method** and is described in Sec. 3.2.

The input data provided to the scoring engine must conform to the **input** type in the sense that there must be an unambiguous way to generate it from Avro-encoded data, though this conversion need not actually take place. For example, if the **input** is **{"type": "array", "items": "int"}**, then the values passed to the scoring engine must be ordered lists of integers, though they may be implemented as arrays, linked lists, immutable vectors, or any other functionally equivalent data structure that the PFA implementation is capable of using in calculations. The data source need not be Avro-encoded; the Avro schema is only used to specify the type, not to perform conversions. Similarly, the output data must conform to the **output** type in the sense that there must be an unambiguous way to convert it to Avro-encoded data, though this conversion need not actually take place.

Given that the input and output types are described by Avro schemae, the Avro binary and JSON data formats would be particularly convenient ways to read and write data. However, there is no requirement that a PFA system should have this capability. Data conversion and internal data format are both outside the scope of the PFA specification.

3.4 Persistent state: cells and pools

PFA defines two mechanisms to maintain state: cells and pools. Cells are global variables with a fixed name and type that must be initialized before the scoring engine's run begins. A cell's value can change to a new value of the same type, but the cell cannot be deleted and new cells cannot be created during the scoring engine's run. Pools are global namespaces with a fixed type. New named values can be created within a pool at runtime, as long as they have the correct type, and old values can be deleted at runtime.

A pool of type "`X`" would be equivalent to a cell whose type is "map of `X`" execpt for performance and concurrency issues. All special forms and library functions in the PFA specification treat data structures as immutable objects (see Sec. 5.1), but scoring engines often need to maintain very large key-value tables. If pools were not available, a PFA implementation would either incur a performance penalty if it maintained a large map in a cell as an immutable object or if it maintained all temporary variables as mutable objects. With both cells and pools, a PFA implementation may maintain all values as immutable objects, including cells, and maintain pools as mutable maps of immutable objects. See Sec. 3.5 for a discussion of concurrency issues in cells and pools.

Cells can only be accessed through the "cell" special form and can only be modified through the "cell-to" special form. Pools can only be accessed through the "pool" special form and modified through the "pool-to" special form (see Sec. 7.7).

One common use of persistent state is to represent a complex statistical model, such as a large decision tree. In most cases, such a model is constant during the scoring engine's run, and this constraint may be enforced through static analysis if the model is stored in a cell. Another common use is to represent recent or accumulated data in a table, indexed by key. In most cases, this table is updated frequently and new table entries may be added at any time. Furthermore, it is often useful to distribute the table-fill operation among a battery of concurrent scoring engines, with different engines modifying different keys at the same time. These cases are more easily implemented as pools or shared pools.

3.5 Concurrent access to shared state

If the `shared` member of a cell or pool's specification is `true`, the cell or pool is not assumed to be thread-local (see Sec. 2.1). It may be shared among a battery of identical scoring engines in a multi-threaded process, shared among identical scoring engines distributed across a network, shared in a database among different types of processes, or shared among components of an integrated circuit, etc. In any case, some rules must be followed to avoid simultaneous attempts to modify the data, and these rules must be standardized to ensure that the same scoring engine has the same behavior on different systems.

Shared cells and pools in PFA follow a read-copy-update rule for concurrent access: attempts to read the shared resource (through the "cell" or "pool" special forms) always succeed without blocking and attempts to write (through the "cell-to" or "pool-to" special forms) lock the resource or wait until another writer's lock is released. The writers must operate on a copy of the cell or pool's data (or on immutable data) so that readers can access the old version during the update process. The new value must be updated atomically at the end of the update process.

Although an update operation may only modify a part of the cell's structure (one value in an array, for instance), the granularity of the writers' lock is the entire cell: two writers must not be able to modify different parts of the same cell at the same time. The granularity of the writers' lock on pools is limited to a single named entity within the pool: two writers must be able to modify different entities in the pool at the same time, but not different parts of the same named entity.

The "cell-to" and "pool-to" special forms accept user-defined update functions (see Sec. 7.7) but these functions must not directly or indirectly call "cell-to" or "pool-to" because such a situation could lead to deadlock. This rule can be enforced by examining the call graph.

3.6 Exceptions

As much as is reasonably possible, PFA documents can be statically analyzed to avoid errors at runtime. However, some error states cannot be predicted without runtime information. These error states and their exact messages are explicitly defined for each susceptible special form and library function. If the specified error conditions are met in the **begin** routine of a scoring engine, processing stops and should not continue to the **action** routine. If error conditions occur when the **action** routine is processing an input datum, processing of that datum stops and may either continue to the next datum or stop the scoring engine entirely. The PFA host may choose to stop or continue on the basis of the standardized error message. If error conditions occur in the **end** routine, processing stops.

This abrupt end of processing may occur deep in an expression or array of expressions and behaves like an exception: control flow exits the routine immediately upon encountering the error condition and is either caught by the host system or it halts the process. In environments where this is difficult to implement, control flow may continue to the end of the routine, but side-effects such as modifications to persistent state and log messages must be avoided.

If the host system catches an exception in **action** and continues to the next datum, and if a cell or pool's **rollback** member is **true**, then that cell or pool should be reverted to the value that it had at the beginning of the **action** (see Sec. 2.1). If the **rollback** member is absent or **false**, then the cell or pool's value at the start of the next **action** should be the value it had at the time of the exception.

In addition to exceptions raised by special forms and library functions, a PFA document can raise custom exceptions with the "error" special form (see Sec. 7.14). The rules described above apply equally to custom exceptions, though we recommend that PFA systems differentiate between built-in exceptions (whose error messages are explicitly defined by this specification) and custom exceptions (whose error messages are free-form).

If a **timeout** is defined in the PFA document's **options** or is imposed by the PFA system, a **begin**, **action**, or **end** routine that exceeds this timeout raises an exception with the message "exceeded timeout of N milliseconds" where N is the relevant timeout. Timeout exceptions follow the same rules as built-in and custom exceptions.

If possible in a given system, no exceptions other than PFA exceptions should ever be raised while executing a **begin**, **action**, or **end** routine.

3.7 Execution options

The **options** top-level field allows PFA documents to request that they are executed in a particular way. However, the PFA system may override any of these options with its own values, or with the defaults. When overriding an option, the PFA system should somehow indicate that this is the case, possibly through a log message.

The complete set of options, their JSON types, and their default values are given below. If a PFA document attempts to set an option that is not in this list or attempts to set an option with the wrong type, it is a semantic error (phase 2 in see Sec. 3.1) and the scoring engine should not be started.

Option name	JSON type	Default value	Description
timeout	integer	−1	Number of milliseconds to let the **begin**, **action**, or **end** routine run; at or after this time, the PFA system may stop the routine with an exception (see Sec. 3.6). If negative, the execution has no timeout.

Option name	JSON type	Default value	Description
`timeout.begin`	integer	`timeout`	A specific timeout for the `begin` routine that overrides the general `timeout`.
`timeout.action`	integer	`timeout`	A specific timeout for the `action` routine that overrides the general `timeout`.
`timeout.end`	integer	`timeout`	A specific timeout for the `end` routine that overrides the general `timeout`.

3.8 Pseudorandom number management

The `randseed` top-level field specifies a seed for library functions that generate pseudorandom numbers. If the `randseed` is absent, the random number generator should be unpredictable: multiple runs of the same PFA document would yield different results if the output depends on pseudorandom numbers. If a `randseed` is provided, the random number generator should be predictable: multiple runs of the same PFA document would yield the same results on the same system. Explicitly setting a `randseed` is useful for tests.

The pseudorandom number generator maintains state between `begin`, `action`, and `end` invocations: the generator is not reseeded with each call. If a PFA document is used to create a battery of identical scoring engines, the `randseed` is used to generate different seeds for all of the scoring engines: they are not guaranteed to produce identical results.

The algorithm for generating pseudorandom numbers is not specified, so different PFA implementations may use different algorithms. Therefore, a PFA document with an explicit `randseed` is only guaranteed to yield identical results when re-run on the same system. On different systems, it may yield different results.

Every library function that depends on pseudorandom numbers should be seeded by the `randseed`. Pseudorandom functions are explicitly denoted by this specification.

4 Type system

4.1 Avro types

Orderless Avro schema reading (ForwardDeclarationParser)

Type-safe null

Limitations: (1) String-only map keys, (2) No "set" object (use array's set-like functions or a map from string values to nulls), (3) No circular references (make a map of string-valued keys: acts as a weak reference)

4.2 Type inference

inputs (`input`, `tally`, function parameters, inline literals, inline new arrays/maps/records, and cell/pool references) are taken as givens. outputs (`output`, function arguments) are derived from the tree of expressions and checked against the declared value.

4.3 Type resolution, promotion, and covariance

copy Avro rules here (and check them in REPL)

4.4 Function parameter patterns

including wildcards and wildrecords; typeset them the same way as they are used in libfcn reference

solution to an equation

promotes conflicting label matches to a union

5 Symbols, scope, and data structures

implicit garbage collector; no restriction on the choice of garbage collector (use whatever is available on your system!)

5.1 Immutable data, reassignable symbols

5.2 Expression-level scope and mutation restrictions

6 User-defined functions

6.1 Syntax and scope

6.2 Anonymous callbacks and function references

7 Expressions

Special forms and ordinary function calls

7.1 Function calls

7.2 Symbol references

7.3 Literal values

7.4 Creating arrays, maps, and records

7.5 Symbol assignment and reassignment

7.6 Extracting from and updating arrays, maps, and records

7.7 Extracting from and updating cells and pools

7.8 Do blocks

7.9 Conditionals: if and cond

7.10 While loops: pretest and posttest

7.11 For loops: by index, array element, and key-value

7.12 Type-safe casting

7.13 Inline documentation

7.14 User-defined exceptions

7.15 Log messages

8 Core library

8.1 Basic arithmetic

8.1.1 Addition of two values (+)

Signature: `{"+": [x, y]}`

`x`	any `A` of {int, long, float, double}
`y`	`A`
(returns)	`A`

Description: Add `x` and `y`.

Details:

Float and double overflows do not produce runtime errors but result in positive or negative infinity, which would be carried through any subsequent calculations (see IEEE 754). Use `impute.ensureFinite` to produce errors from infinite or NaN values.

Runtime Errors:

Integer results above or below -2147483648 and 2147483647 (inclusive) produce an "int overflow" runtime error.

Long-integer results above or below -9223372036854775808 and 9223372036854775807 (inclusive) produce a "long overflow" runtime error.

8.1.2 Subtraction (−)

Signature: `{"-": [x, y]}`

`x`	any `A` of {int, long, float, double}
`y`	`A`
(returns)	`A`

Description: Subtract `y` from `x`.

Details:

Float and double overflows do not produce runtime errors but result in positive or negative infinity, which would be carried through any subsequent calculations (see IEEE 754). Use `impute.ensureFinite` to produce errors from infinite or NaN values.

Runtime Errors:

Integer results above or below -2147483648 and 2147483647 (inclusive) produce an "int overflow" runtime error.

Long-integer results above or below -9223372036854775808 and 9223372036854775807 (inclusive) produce a "long overflow" runtime error.

8.1.3 Multiplication of two values (*)

Signature: `{"*": [x, y]}`

`x`	any `A` of {int, long, float, double}
`y`	`A`
(returns)	`A`

Description: Multiply `x` and `y`.

Details:

Float and double overflows do not produce runtime errors but result in positive or negative infinity, which would be carried through any subsequent calculations (see IEEE 754). Use `impute.ensureFinite` to produce errors from infinite or NaN values.

Runtime Errors:

Integer results above or below -2147483648 and 2147483647 (inclusive) produce an "int overflow" runtime error.

Long-integer results above or below -9223372036854775808 and 9223372036854775807 (inclusive) produce a "long overflow" runtime error.

8.1.4 Floating-point division (/)

Signature: `{"/": [x, y]}`

`x`	double
`y`	double
(returns)	double

Description: Divide `y` from `x`, returning a floating-point number (even if `x` and `y` are integers).

8.1.5 Integer division (//)

Signature: `{"//": [x, y]}`

`x`	any `A` of {int, long}
`y`	`A`
(returns)	`A`

Description: Divide `y` from `x`, returning the largest whole number `N` for which `N` $\leq$ `x`/`y` (integral floor division).

8.1.6 Negation (u−)

Signature: `{"u-": [x]}`

`x` any `A` of {int, long, float, double}

(returns) `A`

Description: Return the additive inverse of `x`.

Runtime Errors:

For exactly one integer value, -2147483648, this function produces an "int overflow" runtime error.

For exactly one long value, -9223372036854775808, this function produces a "long overflow" runtime error.

8.1.7 Modulo (%)

Signature: `{"%": [k, n]}`

`k` any `A` of {int, long, float, double}

`n` `A`

(returns) `A`

Description: Return `k` modulo `n`; the result has the same sign as the modulus `n`.

Details:

This is the behavior of the `%` operator in Python, `mod`/`modulo` in Ada, Haskell, and Scheme.

8.1.8 Remainder (%%)

Signature: `{"%%": [k, n]}`

`k` any `A` of {int, long, float, double}

`n` `A`

(returns) `A`

Description: Return the remainder of `k` divided by `n`; the result has the same sign as the dividend `k`.

Details:

This is the behavior of the `%` operator in Fortran, C/C++, and Java, `rem`/`remainder` in Ada, Haskell, and Scheme.

8.1.9 Raising to a power (**)

Signature: `{"**": [x, y]}`

`x` any `A` of {int, long, float, double}

`y` `A`

(returns) `A`

Description: Raise `x` to the power `n`.

Details:

Float and double overflows do not produce runtime errors but result in positive or negative infinity, which would be carried through any subsequent calculations (see IEEE 754). Use `impute.ensureFinite` to produce errors from infinite or NaN values.

Runtime Errors:

Integer results above or below -2147483648 and 2147483647 (inclusive) produce an "int overflow" runtime error.

Long-integer results above or below -9223372036854775808 and 9223372036854775807 (inclusive) produce a "long overflow" runtime error.

8.2 Comparison operators

Avro defines a sort order for every pair of values with a compatible type, so any two objects of compatible type can be compared in PFA.

8.2.1 General comparision (cmp)

Signature: `{"cmp": [x, y]}`

`x`	any `A`
`y`	`A`
(returns)	int

Description: Return `1` if `x` is greater than `y`, `-1` if `x` is less than `y`, and `0` if `x` and `y` are equal.

8.2.2 Equality (==)

Signature: `{"==": [x, y]}`

`x`	any `A`
`y`	`A`
(returns)	boolean

Description: Return `true` if `x` is equal to `y`, `false` otherwise.

8.2.3 Inequality (!=)

Signature: `{"!=": [x, y]}`

`x`	any `A`
`y`	`A`
(returns)	boolean

Description: Return `true` if `x` is not equal to `y`, `false` otherwise.

8.2.4 Less than (<)

Signature: `{"<": [x, y]}`

`x`	any `A`
`y`	`A`
(returns)	boolean

Description: Return `true` if `x` is less than `y`, `false` otherwise.

8.2.5 Less than or equal to (<=)

Signature: `{"<=": [x, y]}`

`x`	any `A`
`y`	`A`
(returns)	boolean

Description: Return `true` if `x` is less than or equal to `y`, `false` otherwise.

8.2.6 Greater than (>)

Signature: `{">": [x, y]}`

`x`	any `A`
`y`	`A`
(returns)	boolean

Description: Return `true` if `x` is greater than `y`, `false` otherwise.

8.2.7 Greater than or equal to (>=)

Signature: `{">=": [x, y]}`

`x`	any `A`
`y`	`A`
(returns)	boolean

Description: Return `true` if `x` is greater than or equal to `y`, `false` otherwise.

8.2.8 Maximum of two values (max)

Signature: `{"max": [x, y]}`

`x`	any `A`
`y`	`A`
(returns)	`A`

Description: Return `x` if `x` $\geq$ `y`, `y` otherwise.

Details:

For the maximum of more than two values, see `a.max`

8.2.9 Minimum of two values (min)

Signature: `{"min": [x, y]}`

`x`	any `A`
`y`	`A`
(returns)	`A`

Description: Return `x` if `x` $<$ `y`, `y` otherwise.

Details:

For the minimum of more than two values, see `a.min`

8.3 Logical operators

8.3.1 Logical and (and)

Signature: `{"and": [x, y]}`

`x`	boolean
`y`	boolean
(returns)	boolean

Description: Return `true` if `x` and `y` are both `true`, `false` otherwise.

Details:

If `x` is `false`, `y` won't be evaluated. (Only relevant for arguments with side effects.)

8.3.2 Logical or (or)

Signature: `{"or": [x, y]}`

`x` boolean

`y` boolean

(returns) boolean

Description: Return `true` if either `x` or `y` (or both) are `true`, `false` otherwise.

Details:

If `x` is `true`, `y` won't be evaluated. (Only relevant for arguments with side effects.)

8.3.3 Logical xor (xor)

Signature: `{"xor": [x, y]}`

`x` boolean

`y` boolean

(returns) boolean

Description: Return `true` if `x` is `true` and `y` is `false` or if `x` is `false` and `y` is `true`, but return `false` for any other case.

8.3.4 Logical not (not)

Signature: `{"not": [x]}`

`x` boolean

(returns) boolean

Description: Return `true` if `x` is `false` and `false` if `x` is `true`.

8.4 Bitwise arithmetic

8.4.1 Bitwise and (&)

Signature: `{"&": [x, y]}`

`x` int

`y` int

(returns) int

or

`x` long

`y` long

(returns) long

Description: Calculate the bitwise-and of `x` and `y`.

8.4.2 Bitwise or (|)

Signature: `{"|": [x, y]}`

`x`	int
`y`	int
(returns)	int
or	
`x`	long
`y`	long
(returns)	long

Description: Calculate the bitwise-or of `x` and `y`.

8.4.3 Bitwise xor (^)

Signature: `{"^": [x, y]}`

`x`	int
`y`	int
(returns)	int
or	
`x`	long
`y`	long
(returns)	long

Description: Calculate the bitwise-exclusive-or of `x` and `y`.

8.4.4 Bitwise not (~)

Signature: `{"~": [x]}`

`x`	int
(returns)	int
or	
`x`	long
(returns)	long

Description: Calculate the bitwise-not of `x`.

9 Math library

9.1 Constants

Constants such as π and e are represented as stateless functions with no arguments. Specific implementations may choose to replace the function call with its inline value.

9.1.1 Archimedes' constant π (m.pi)

Signature: `{"m.pi": []}`

(returns) double

Description: The double-precision number that is closer than any other to π, the ratio of a circumference of a circle to its diameter.

9.1.2 Euler's constant e (m.e)

Signature: `{"m.e": []}`

(returns) double

Description: The double-precision number that is closer than any other to e, the base of natural logarithms.

9.2 Common functions

9.2.1 Square root (m.sqrt)

Signature: `{"m.sqrt": [x]}`

`x` double

(returns) double

Description: Return the positive square root of `x`.

Details:

The domain of this function is from 0 (inclusive) to infinity. Beyond this domain, the result is Use

9.2.2 Hypotnuse (m.hypot)

Signature: `{"m.hypot": [x, y]}`

`x` double

`y` double

(returns) double

Description: Return $\sqrt{x^2 + y^2}$.

Details:

Avoids round-off or overflow errors in the intermediate steps.

The domain of this function is the whole real line; no input is invalid.

9.2.3 Trigonometric sine (m.sin)

Signature: `{"m.sin": [x]}`

`x` double

(returns) double

Description: Return the trigonometric sine of `x`, which is assumed to be in radians.

Details:

The domain of this function is the whole real line; no input is invalid.

9.2.4 Trigonometric cosine (m.cos)

Signature: `{"m.cos": [x]}`

`x` double

(returns) double

Description: Return the trigonometric cosine of `x`, which is assumed to be in radians.

Details:

The domain of this function is the whole real line; no input is invalid.

9.2.5 Trigonometric tangent (m.tan)

Signature: `{"m.tan": [x]}`

`x` double

(returns) double

Description: Return the trigonometric tangent of `x`, which is assumed to be in radians.

Details:

The domain of this function is the whole real line; no input is invalid.

9.2.6 Inverse trigonometric sine (m.asin)

Signature: `{"m.asin": [x]}`

`x`	double
(returns)	double

Description: Return the arc-sine (inverse of the sine function) of `x` as an angle in radians between $-\pi/2$ and $\pi/2$.

Details:

The domain of this function is from -1 to 1 (inclusive). Beyond this domain, the result is Use

9.2.7 Inverse trigonometric cosine (m.acos)

Signature: `{"m.acos": [x]}`

`x`	double
(returns)	double

Description: Return the arc-cosine (inverse of the cosine function) of `x` as an angle in radians between 0 and π.

Details:

The domain of this function is from -1 to 1 (inclusive). Beyond this domain, the result is Use

9.2.8 Inverse trigonometric tangent (m.atan)

Signature: `{"m.atan": [x]}`

`x`	double
(returns)	double

Description: Return the arc-tangent (inverse of the tangent function) of `x` as an angle in radians between $-\pi/2$ and $\pi/2$.

Details:

The domain of this function is the whole real line; no input is invalid.

9.2.9 Robust inverse trigonometric tangent (m.atan2)

Signature: `{"m.atan2": [y, x]}`

`y` double

`x` double

(returns) double

Description: Return the arc-tangent (inverse of the tangent function) of `y`/`x` without loss of precision for small `x`.

Details:

The domain of this function is the whole real plane; no pair of inputs is invalid.

Note that `y` is the first parameter and `x` is the second parameter.

9.2.10 Hyperbolic sine (m.sinh)

Signature: `{"m.sinh": [x]}`

`x` double

(returns) double

Description: Return the hyperbolic sine of `x`, which is equal to $\frac{e^x - e^{-x}}{2}$.

Details:

The domain of this function is the whole real line; no input is invalid.

9.2.11 Hyperbolic cosine (m.cosh)

Signature: `{"m.cosh": [x]}`

`x` double

(returns) double

Description: Return the hyperbolic cosine of `x`, which is equal to $\frac{e^x + e^{-x}}{2}$

Details:

The domain of this function is the whole real line; no input is invalid.

9.2.12 Hyperbolic tangent (m.tanh)

Signature: `{"m.tanh": [x]}`

`x` double

(returns) double

Description: Return the hyperbolic tangent of `x`, which is equal to $\frac{e^x - e^{-x}}{e^x + e^{-x}}$.

Details:

The domain of this function is the whole real line; no input is invalid.

9.2.13 Natural exponential (m.exp)

Signature: `{"m.exp": [x]}`

`x`	double
(returns)	double

Description: Return `m.e` raised to the power of `x`.

Details:

The domain of this function is the whole real line; no input is invalid.

9.2.14 Natural exponential minus one (m.expm1)

Signature: `{"m.expm1": [x]}`

`x`	double
(returns)	double

Description: Return $e^x - 1$.

Details:

Avoids round-off or overflow errors in the intermediate steps.

The domain of this function is the whole real line; no input is invalid.

9.2.15 Natural logarithm (m.ln)

Signature: `{"m.ln": [x]}`

`x`	double
(returns)	double

Description: Return the natural logarithm of `x`.

Details:

The domain of this function is from 0 to infinity (exclusive). Given zero, the result is negative infinity, and below zero, the result is Use

9.2.16 Logarithm base 10 (m.log10)

Signature: `{"m.log10": [x]}`

`x` double

(returns) double

Description: Return the logarithm base 10 of `x`.

Details:

The domain of this function is from 0 to infinity (exclusive). Given zero, the result is negative infinity, and below zero, the result is Use

9.2.17 Arbitrary logarithm (m.log)

Signature: `{"m.log": [x, base]}`

`x` double

`base` int

(returns) double

Description: Return the logarithm of `x` with a given `base`.

Details:

The domain of this function is from 0 to infinity (exclusive). Given zero, the result is negative infinity, and below zero, the result is Use

Runtime Errors:

If `base` is less than or equal to zero, this function produces a "base must be positive" runtime error.

9.2.18 Natural logarithm of one plus square (m.ln1p)

Signature: `{"m.ln1p": [x]}`

`x` double

(returns) double

Description: Return $ln(x^2 + 1)$.

Details:

Avoids round-off or overflow errors in the intermediate steps.

The domain of this function is from -1 to infinity (exclusive). Given -1, the result is negative infinity, and below -1, the result is Use

9.3 Rounding

9.3.1 Absolute value (m.abs)

Signature: `{"m.abs": [x]}`

`x` any `A` of {int, long, float, double}

(returns) `A`

Description: Return the absolute value of `x`.

Details:

The domain of this function is the whole real line; no input is invalid.

Runtime Errors:

For exactly one integer value, -2147483648, this function produces an "int overflow" runtime error.

For exactly one long value, -9223372036854775808, this function produces a "long overflow" runtime error.

9.3.2 Floor (m.floor)

Signature: `{"m.floor": [x]}`

`x` double

(returns) double

Description: Return the largest (closest to positive infinity) whole number that is less than or equal to the input.

Details:

The domain of this function is the whole real line; no input is invalid.

9.3.3 Ceiling (m.ceil)

Signature: `{"m.ceil": [x]}`

`x` double

(returns) double

Description: Return the smallest (closest to negative infinity, not closest to zero) whole number that is greater than or equal to the input.

Details:

The domain of this function is the whole real line; no input is invalid.

9.3.4 Simple rounding (m.round)

Signature: `{"m.round": [x]}`

`x` float

(returns) int

or

`x`	double
(returns)	long

Description: Return the closest whole number to `x`, rounding up if the fractional part is exactly one-half.

Details:

Equal to `m.floor` of (`x` + 0.5).

Runtime Errors:

Integer results outside of -2147483648 and 2147483647 (inclusive) produce an "int overflow" runtime error.

Long-integer results outside of -9223372036854775808 and 9223372036854775807 (inclusive) produce a "long overflow" runtime error.

9.3.5 Unbiased rounding (m.rint)

Signature: `{"m.rint": [x]}`

`x`	double
(returns)	double

Description: Return the closest whole number to `x`, rounding toward the nearest even number if the fractional part is exactly one-half.

9.3.6 Threshold function (m.signum)

Signature: `{"m.signum": [x]}`

`x`	double
(returns)	int

Description: Return 0 if `x` is zero, 1 if `x` is positive, and -1 if `x` is negative.

Details:

The domain of this function is the whole real line; no input is invalid.

9.3.7 Copy sign (m.copysign)

Signature: `{"m.copysign": [mag, sign]}`

`mag`	any `A` of {int, long, float, double}
`sign`	`A`
(returns)	`A`

Description: Return a number with the magnitude of `mag` and the sign of `sign`.

Details:

The domain of this function is the whole real or integer plane; no pair of inputs is invalid.

9.4 Linear algebra

including named row/col matrices

10 String manipulation

Strings are immutable, so none of the following functions modifies a string in-place. Some return a modified version of the original string.

10.1 Basic access

10.1.1 Length (s.len)

Signature: `{"s.len": [s]}`

`s`	string
(returns)	int

Description: Return the length of string `s`.

10.1.2 Extract substring (s.substr)

Signature: `{"s.substr": [s, start, end]}`

`s`	string
`start`	int
`end`	int
(returns)	string

Description: Return the substring of `s` from `start` (inclusive) until `end` (exclusive).

Details:

Negative indexes count from the right (-1 is just before the last item), indexes beyond the legal range are truncated, and `end` $\leq$ `start` specifies a zero-length subsequence just before the `start` character. All of these rules follow Python's slice behavior.

10.1.3 Modify substring (s.substrto)

Signature: `{"s.substrto": [s, start, end, replacement]}`

`s`	string
`start`	int
`end`	int
`replacement`	string
(returns)	string

Description: Replace `s` from `start` (inclusive) until `end` (exclusive) with `replacement`.

Details:

Negative indexes count from the right (-1 is just before the last item), indexes beyond the legal range are truncated, and **`end`** $\leq$ **`start`** specifies a zero-length subsequence just before the **`start`** character. All of these rules follow Python's slice behavior.

10.2 Search and replace

10.2.1 Contains (s.contains)

Signature: `{"s.contains": [haystack, needle]}`

`haystack`	string
`needle`	string
(returns)	boolean

Description: Return `true` if `haystack` contains `needle`, `false` otherwise.

10.2.2 Count instances (s.count)

Signature: `{"s.count": [haystack, needle]}`

`haystack`	string
`needle`	string
(returns)	int

Description: Count the number of times `needle` appears in `haystack`.

10.2.3 Find first index (s.index)

Signature: `{"s.index": [haystack, needle]}`

`haystack`	string
`needle`	string
(returns)	int

Description: Return the lowest index where `haystack` contains `needle` or -1 if `haystack` does not contain `needle`.

10.2.4 Find last index (s.rindex)

Signature: `{"s.rindex": [haystack, needle]}`

`haystack`	string
`needle`	string
(returns)	int

Description: Return the highest index where `haystack` contains `needle` or -1 if `haystack` does not contain `needle`.

10.2.5 Check start (s.startswith)

Signature: `{"s.startswith": [haystack, needle]}`

`haystack`	string
`needle`	string
(returns)	boolean

Description: Return `true` if the first (leftmost) subseqence of `haystack` is equal to `needle`, false otherwise.

10.2.6 Check end (s.endswith)

Signature: `{"s.endswith": [haystack, needle]}`

`haystack`	string
`needle`	string
(returns)	boolean

Description: Return `true` if the last (rightmost) subseqence of `haystack` is equal to `needle`, false otherwise.

10.3 Conversions to or from other types

10.3.1 Join an array of strings (s.join)

Signature: `{"s.join": [array, sep]}`

`array`	array of string
`sep`	string
(returns)	string

Description: Combine strings from `array` into a single string, delimited by `sep`.

10.3.2 Split into an array of strings (s.split)

Signature: `{"s.split": [s, sep]}`

`s`	string
`sep`	string
(returns)	array of string

Description: Divide a string into an array of substrings, splitting at and removing delimiters `sep`.

Details:

If `s` does not contain `sep`, this function returns an array whose only element is `s`. If `sep` appears at the beginning or end of `s`, the array begins with or ends with an empty string. These conventions match Python's behavior.

10.4 Conversions to or from other strings

10.4.1 Concatenate two strings (s.concat)

Signature: `{"s.concat": [x, y]}`

`x`	string
`y`	string
(returns)	string

Description: Append `y` to `x` to form a single string.

Details:

To concatenate an array of strings, use s.join with an empty string as `sep`.

10.4.2 Repeat pattern (s.repeat)

Signature: `{"s.repeat": [s, n]}`

`s`	string
`n`	int
(returns)	string

Description: Create a string by concatenating `s` with itself `n` times.

10.4.3 Lowercase (s.lower)

Signature: `{"s.lower": [s]}`

`s` string

(returns) string

Description: Convert `s` to lower-case.

10.4.4 Uppercase (s.upper)

Signature: `{"s.upper": [s]}`

`s` string

(returns) string

Description: Convert `s` to upper-case.

10.4.5 Left-strip (s.lstrip)

Signature: `{"s.lstrip": [s, chars]}`

`s` string

`chars` string

(returns) string

Description: Remove any characters found in `chars` from the beginning (left) of `s`.

Details:

The order of characters in `chars` is irrelevant.

10.4.6 Right-strip (s.rstrip)

Signature: `{"s.rstrip": [s, chars]}`

`s` string

`chars` string

(returns) string

Description: Remove any characters found in `chars` from the end (right) of `s`.

Details:

The order of characters in `chars` is irrelevant.

10.4.7 Strip both ends (s.strip)

Signature: `{"s.strip": [s, chars]}`

`s`	string
`chars`	string
(returns)	string

Description: Remove any characters found in `chars` from the beginning or end of `s`.

Details:

The order of characters in `chars` is irrelevant.

10.4.8 Replace all matches (s.replaceall)

Signature: `{"s.replaceall": [s, original, replacement]}`

`s`	string
`original`	string
`replacement`	string
(returns)	string

Description: Replace every instance of the substring `original` from `s` with `replacement`.

10.4.9 Replace first match (s.replacefirst)

Signature: `{"s.replacefirst": [s, original, replacement]}`

`s`	string
`original`	string
`replacement`	string
(returns)	string

Description: Replace the first (leftmost) instance of the substring `original` from `s` with `replacement`.

10.4.10 Replace last match (s.replacelast)

Signature: `{"s.replacelast": [s, original, replacement]}`

`s`	string
`original`	string
`replacement`	string
(returns)	string

Description: Replace the last (rightmost) instance of the substring `original` from `s` with `replacement`.

10.4.11 Translate characters (s.translate)

Signature: `{"s.translate": [s, oldchars, newchars]}`

`s`	string
`oldchars`	string
`newchars`	string
(returns)	string

Description: For each character in `s` that is also in `oldchars` with some index `i`, replace it with the character at index `i` in `newchars`. Any character in `s` that is not in `oldchars` is unchanged. Any index `i` that is greater than the length of `newchars` is replaced with nothing.

Details:

This is the behavior of the the Posix command `tr`, where `s` takes the place of standard input and `oldchars` and `newchars` are the `tr` commandline options.

10.5 Regular Expressions

and stemming

11 Array Manipulation

11.1 Basic access

11.1.1 Length (a.len)

Signature: `{"a.len": [a]}`

`a`	array of any `A`
(returns)	int

Description: Return the length of array `a`.

11.1.2 Extract subsequence (a.subseq)

Signature: `{"a.subseq": [a, start, end]}`

`a`	array of any `A`
`start`	int
`end`	int
(returns)	array of `A`

Description: Return the subsequence of `a` from `start` (inclusive) until `end` (exclusive).

Details:

Negative indexes count from the right (-1 is just before the last item), indexes beyond the legal range are truncated, and `end` $\leq$ `start` specifies a zero-length subsequence just before the `start` character. All of these rules follow Python's slice behavior.

11.1.3 Modify subsequence (a.subseqto)

Signature: `{"a.subseqto": [a, start, end, replacement]}`

`a`	array of any `A`
`start`	int
`end`	int
`replacement`	array of `A`
(returns)	array of `A`

Description: Return a new array by replacing `a` from `start` (inclusive) until `end` (exclusive) with `replacement`.

Details:

Negative indexes count from the right (-1 is just before the last item), indexes beyond the legal range are truncated, and **`end`** $\leq$ **`start`** specifies a zero-length subsequence just before the **`start`** character. All of these rules follow Python's slice behavior.

Note: **`a`** is not changed in-place; this is a side-effect-free function.

11.2 Search and replace

11.2.1 Contains (a.contains)

Signature: **`{"a.contains": [haystack, needle]}`**

`haystack`	array of any **`A`**
`needle`	array of **`A`**
(returns)	boolean
or	
`haystack`	array of any **`A`**
`needle`	**`A`**
(returns)	boolean

Description: Return **`true`** if **`haystack`** contains **`needle`**, **`false`** otherwise.

11.2.2 Count instances (a.count)

Signature: **`{"a.count": [haystack, needle]}`**

`haystack`	array of any **`A`**
`needle`	array of **`A`**
(returns)	int
or	
`haystack`	array of any **`A`**
`needle`	**`A`**
(returns)	int

Description: Count the number of times **`needle`** appears in **`haystack`**.

11.2.3 Count instances by predicate (a.countPredicate)

libfcna.countPredicate

11.2.4 Find first index (a.index)

Signature: **`{"a.index": [haystack, needle]}`**

`haystack`	array of any `A`
`needle`	array of `A`
(returns)	int
or	
`haystack`	array of any `A`
`needle`	`A`
(returns)	int

Description: Return the lowest index where `haystack` contains `needle` or -1 if `haystack` does not contain `needle`.

11.2.5 Find last index (a.rindex)

Signature: `{"a.rindex": [haystack, needle]}`

`haystack`	array of any `A`
`needle`	array of `A`
(returns)	int
or	
`haystack`	array of any `A`
`needle`	`A`
(returns)	int

Description: Return the highest index where `haystack` contains `needle` or -1 if `haystack` does not contain `needle`.

11.2.6 Check start (a.startswith)

Signature: `{"a.startswith": [haystack, needle]}`

`haystack`	array of any `A`
`needle`	array of `A`
(returns)	boolean
or	
`haystack`	array of any `A`
`needle`	`A`
(returns)	boolean

Description: Return `true` if the first (leftmost) subseqence of `haystack` is equal to `needle`, false otherwise.

11.2.7 Check end (a.endswith)

Signature: `{"a.endswith": [haystack, needle]}`

`haystack`	array of any `A`
`needle`	array of `A`
(returns)	boolean
or	
`haystack`	array of any `A`
`needle`	`A`
(returns)	boolean

Description: Return `true` if the last (rightmost) subseqence of `haystack` is equal to `needle`, false otherwise.

11.3 Manipulation

11.3.1 Concatenate two arrays (a.concat)

Signature: `{"a.concat": [a, b]}`

`a`	array of any `A`
`b`	array of `A`
(returns)	array of `A`

Description: Concatenate `a` and `b` to make a new array of the same type.

Details:

The length of the returned array is the sum of the lengths of `a` and `b`.

11.3.2 Append (a.append)

Signature: `{"a.append": [a, item]}`

`a`	array of any `A`
`item`	`A`
(returns)	array of `A`

Description: Return a new array by adding `item` at the end of `a`.

Details:

Note: `a` is not changed in-place; this is a side-effect-free function.

The length of the returned array is one more than `a`.

11.3.3 Insert or prepend (a.insert)

Signature: `{"a.insert": [a, index, item]}`

`a`	array of any `A`
`index`	int
`item`	`A`
(returns)	array of `A`

Description: Return a new array by inserting `item` at `index` of `a`.

Details:

Negative indexes count from the right (-1 is just before the last item), following Python's index behavior.

Note: `a` is not changed in-place; this is a side-effect-free function.

The length of the returned array is one more than `a`.

Runtime Errors:

If `index` is beyond the range of `a`, an "array out of range" runtime error is raised.

11.3.4 Replace item (a.replace)

Signature: `{"a.replace": [a, index, item]}`

`a`	array of any `A`
`index`	int
`item`	`A`
(returns)	array of `A`

Description: Return a new array by replacing `index` of `a` with `item`.

Details:

Negative indexes count from the right (-1 is just before the last item), following Python's index behavior.

Note: `a` is not changed in-place; this is a side-effect-free function.

The length of the returned array is equal to that of `a`.

Runtime Errors:

If `index` is beyond the range of `a`, an "array out of range" runtime error is raised.

11.3.5 Remove item (a.remove)

Signature: `{"a.remove": [a, start, end]}` or `{"a.remove": [a, index]}`

`a`	array of any `A`
`start`	int
`end`	int
(returns)	array of `A`

or

`a`	array of any `A`
`index`	int
(returns)	array of `A`

Description: Return a new array by removing elements from `a` from `start` (inclusive) until `end` (exclusive) or just a single `index`.

Details:

Negative indexes count from the right (-1 is just before the last item), indexes beyond the legal range are truncated, and `end` $\leq$ `start` specifies a zero-length subsequence just before the `start` character. All of these rules follow Python's slice behavior.

Note: `a` is not changed in-place; this is a side-effect-free function.

The length of the returned array is one less than `a`.

Runtime Errors:

If `index` is beyond the range of `a`, an "array out of range" runtime error is raised.

11.4 Reordering

11.4.1 Sort (a.sort)

Signature: `{"a.sort": [a]}`

`a`	array of any `A`
(returns)	array of `A`

Description: Return an array with the same elements as `a` but in ascending order (as defined by Avro's sort order).

Details:

Note: `a` is not changed in-place; this is a side-effect-free function.

11.4.2 Sort with a less-than function (a.sortLT)

Signature: `{"a.sortLT": [a, lessThan]}`

`a`	array of any `A`
`lessThan`	function (`A`, `A`) → boolean
(returns)	array of `A`

Description: Return an array with the same elements as `a` but in ascending order as defined by the `lessThan` function.

Details:

Note: `a` is not changed in-place; this is a side-effect-free function.

11.4.3 Randomly shuffle array (a.shuffle)

Signature: `{"a.shuffle": [a]}`

`a`	array of any `A`
(returns)	array of `A`

Description: Return an array with the same elements as `a` but in a random order.

Details:

Note: `a` is not changed in-place; this is a side-effect-free function (except for updating the random number generator).

11.4.4 Reverse order (a.reverse)

Signature: `{"a.reverse": [a]}`

`a`	array of any `A`
(returns)	array of `A`

Description: Return the elements of `a` in reversed order.

11.5 Extreme values

11.5.1 Maximum of all values (a.max)

Signature: `{"a.max": [a]}`

`a`	array of any `A`
(returns)	`A`

Description: Return the maximum value in `a` (as defined by Avro's sort order).

Runtime Errors:

If `a` is empty, an "empty array" runtime error is raised.

11.5.2 Minimum of all values (a.min)

Signature: `{"a.min": [a]}`

`a`	array of any `A`
(returns)	`A`

Description: Return the minimum value in `a` (as defined by Avro's sort order).

Runtime Errors:

If `a` is empty, an "empty array" runtime error is raised.

11.5.3 Maximum with a less-than function (a.maxLT)

Signature: `{"a.maxLT": [a, lessThan]}`

`a`	array of any `A`
`lessThan`	function (`A`, `A`) → boolean
(returns)	`A`

Description: Return the maximum value in `a` as defined by the `lessThan` function.

Runtime Errors:

If `a` is empty, an "empty array" runtime error is raised.

11.5.4 Minimum with a less-than function (a.minLT)

Signature: `{"a.minLT": [a, lessThan]}`

`a`	array of any `A`
`lessThan`	function (`A`, `A`) → boolean
(returns)	`A`

Description: Return the minimum value in `a` as defined by the `lessThan` function.

Runtime Errors:

If `a` is empty, an "empty array" runtime error is raised.

11.5.5 Maximum N items (a.maxN)

Signature: `{"a.maxN": [a, n]}`

`a`	array of any `A`
`n`	int
(returns)	array of `A`

Description: Return the `n` highest values in `a` (as defined by Avro's sort order).

Runtime Errors:

If `a` is empty, an "empty array" runtime error is raised.

If `n` is negative, an "n $<$ 0" runtime error is raised.

11.5.6 Minimum N items (a.minN)

Signature: `{"a.minN": [a, n]}`

`a`	array of any `A`
`n`	int
(returns)	array of `A`

Description: Return the `n` lowest values in `a` (as defined by Avro's sort order).

Runtime Errors:

If `a` is empty, an "empty array" runtime error is raised.

If `n` is negative, an "n $<$ 0" runtime error is raised.

11.5.7 Maximum N with a less-than function (a.maxNLT)

Signature: `{"a.maxNLT": [a, n, lessThan]}`

`a`	array of any `A`
`n`	int
`lessThan`	function (`A`, `A`) $\rightarrow$ boolean
(returns)	array of `A`

Description: Return the `n` highest values in `a` as defined by the `lessThan` function.

Runtime Errors:

If `a` is empty, an "empty array" runtime error is raised.

If `n` is negative, an "n $<$ 0" runtime error is raised.

11.5.8 Minimum N with a less-than function (a.minNLT)

Signature: `{"a.minNLT": [a, n, lessThan]}`

`a`	array of any `A`
`n`	int
`lessThan`	function (`A`, `A`) → boolean
(returns)	array of `A`

Description: Return the `n` lowest values in `a` as defined by the `lessThan` function.

Runtime Errors:

If `a` is empty, an "empty array" runtime error is raised.

If `n` is negative, an "n < 0" runtime error is raised.

11.5.9 Argument maximum (a.argmax)

Signature: `{"a.argmax": [a]}`

`a`	array of any `A`
(returns)	int

Description: Return the index of the maximum value in `a` (as defined by Avro's sort order).

Details:

If the maximum is not unique, this function returns the index of the first maximal value.

Runtime Errors:

If `a` is empty, an "empty array" runtime error is raised.

11.5.10 Argument minimum (a.argmin)

Signature: `{"a.argmin": [a]}`

`a`	array of any `A`
(returns)	int

Description: Return the index of the minimum value in `a` (as defined by Avro's sort order).

Details:

If the minimum is not unique, this function returns the index of the first minimal value.

Runtime Errors:

If `a` is empty, an "empty array" runtime error is raised.

11.5.11 Argument maximum with a less-than function (a.argmaxLT)

Signature: `{"a.argmaxLT": [a, lessThan]}`

`a`	array of any `A`
`lessThan`	function (`A`, `A`) → boolean
(returns)	int

Description: Return the index of the maximum value in `a` as defined by the `lessThan` function.

Details:

If the maximum is not unique, this function returns the index of the first maximal value.

Runtime Errors:

If `a` is empty, an "empty array" runtime error is raised.

11.5.12 Argument minimum with a less-than function (a.argminLT)

Signature: `{"a.argminLT": [a, lessThan]}`

`a`	array of any `A`
`lessThan`	function (`A`, `A`) → boolean
(returns)	int

Description: Return the index of the minimum value in `a` as defined by the `lessThan` function.

Details:

If the minimum is not unique, this function returns the index of the first minimal value.

Runtime Errors:

If `a` is empty, an "empty array" runtime error is raised.

11.5.13 Maximum N arguments (a.argmaxN)

Signature: `{"a.argmaxN": [a, n]}`

`a`	array of any `A`
`n`	int
(returns)	array of int

Description: Return the indexes of the `n` highest values in `a` (as defined by Avro's sort order).

Details:

If any values are not unique, their indexes will be returned in ascending order.

Runtime Errors:

If `a` is empty, an "empty array" runtime error is raised.

If `n` is negative, an "n < 0" runtime error is raised.

11.5.14 Minimum N arguments (a.argminN)

Signature: `{"a.argminN": [a, n]}`

`a`	array of any `A`
`n`	int
(returns)	array of int

Description: Return the indexes of the `n` lowest values in `a` (as defined by Avro's sort order).

Details:

If any values are not unique, their indexes will be returned in ascending order.

Runtime Errors:

If `a` is empty, an "empty array" runtime error is raised.

If `n` is negative, an "n < 0" runtime error is raised.

11.5.15 Maximum N arguments with a less-than function (a.argmaxNLT)

Signature: `{"a.argmaxNLT": [a, n, lessThan]}`

`a`	array of any `A`
`n`	int
`lessThan`	function (`A`, `A`) → boolean
(returns)	array of int

Description: Return the indexes of the `n` highest values in `a` as defined by the `lessThan` function.

Details:

If any values are not unique, their indexes will be returned in ascending order.

Runtime Errors:

If `a` is empty, an "empty array" runtime error is raised.

If `n` is negative, an "n < 0" runtime error is raised.

11.5.16 Minimum N arguments with a less-than function (a.argminNLT)

Signature: `{"a.argminNLT": [a, n, lessThan]}`

`a`	array of any `A`
`n`	int
`lessThan`	function (`A`, `A`) → boolean
(returns)	array of int

Description: Return the indexes of the `n` lowest values in `a` as defined by the `lessThan` function.

Details:

If any values are not unique, their indexes will be returned in ascending order.

Runtime Errors:

If `a` is empty, an "empty array" runtime error is raised.
If `n` is negative, an "n < 0" runtime error is raised.

11.6 Numerical combinations

11.6.1 Add all array values (a.sum)

Signature: `{"a.sum": [a]}`

`a`	array of any `A` of {int, long, float, double}
(returns)	`A`

Description: Return the sum of numbers in `a`.

Details:

Returns zero if the array is empty.

11.6.2 Multiply all array values (a.product)

Signature: `{"a.product": [a]}`

`a`	array of any `A` of {int, long, float, double}
(returns)	`A`

Description: Return the product of numbers in `a`.

Details:

Returns one if the array is empty.

11.6.3 Sum of logarithms (a.lnsum)

Signature: `{"a.lnsum": [a]}`

`a` array of double

(returns) double

Description: Return the sum of the natural logarithm of numbers in `a`.

Details:

Returns zero if the array is empty and `NaN` if any value in the array is zero or negative.

11.6.4 Arithmetic mean (a.mean)

Signature: `{"a.mean": [a]}`

`a` array of double

(returns) double

Description: Return the arithmetic mean of numbers in `a`.

Details:

Returns `NaN` if the array is empty.

11.6.5 Geometric mean (a.geomean)

Signature: `{"a.geomean": [a]}`

`a` array of double

(returns) double

Description: Return the geometric mean of numbers in `a`.

Details:

Returns `NaN` if the array is empty.

11.6.6 Median (a.median)

Signature: `{"a.median": [a]}`

`a` array of any `A`

(returns) `A`

Description: Return the value that is in the center of a sorted version of `a`.

Details:

If `a` has an odd number of elements, the median is the exact center of the sorted array. If `a` has an even number of elements and is a `float` or `double`, the median is the average of the two elements closest to the center of the sorted array. For any other type, the median is the left (first) of the two elements closest to the center of the sorted array.

Runtime Errors:

If `a` is empty, an "empty array" runtime error is raised.

11.6.7 Mode, or most common value (a.mode)

Signature: `{"a.mode": [a]}`

`a`	array of any `A`
(returns)	`A`

Description: Return the mode (most common) value of `a`.

Details:

If several different values are equally common, the median of these is returned.

Runtime Errors:

If `a` is empty, an "empty array" runtime error is raised.

11.7 Set or set-like functions

PFA does not have a set datatype, but arrays can be interpreted as sets with the following functions.

11.7.1 Distinct items (a.distinct)

Signature: `{"a.distinct": [a]}`

`a`	array of any `A`
(returns)	array of `A`

Description: Return an array with the same contents as `a` but with duplicates removed.

11.7.2 Set equality (a.seteq)

Signature: `{"a.seteq": [a, b]}`

`a`	array of any `A`
`b`	array of `A`
(returns)	boolean

Description: Return `true` if `a` and `b` are equivalent, ignoring order and duplicates, `false` otherwise.

11.7.3 Union (a.union)

Signature: `{"a.union": [a, b]}`

`a`	array of any `A`
`b`	array of `A`
(returns)	array of `A`

Description: Return an array that represents the union of `a` and `b`, treated as sets (ignoring order and duplicates).

11.7.4 Intersection (a.intersect)

Signature: `{"a.intersect": [a, b]}`

`a`	array of any `A`
`b`	array of `A`
(returns)	array of `A`

Description: Return an array that represents the intersection of `a` and `b`, treated as sets (ignoring order and duplicates).

11.7.5 Set difference (a.diff)

Signature: `{"a.diff": [a, b]}`

`a`	array of any `A`
`b`	array of `A`
(returns)	array of `A`

Description: Return an array that represents the difference of `a` and `b`, treated as sets (ignoring order and duplicates).

11.7.6 Symmetric set difference (a.symdiff)

Signature: `{"a.symdiff": [a, b]}`

`a`	array of any `A`
`b`	array of `A`
(returns)	array of `A`

Description: Return an array that represents the symmetric difference of `a` and `b`, treated as sets (ignoring order and duplicates).

Details:

The symmetric difference is (`a` diff `b`) union (`b` diff `a`).

11.7.7 Subset check (a.subset)

Signature: `{"a.subset": [little, big]}`

`little`	array of any `A`
`big`	array of `A`
(returns)	boolean

Description: Return `true` if `little` is a subset of `big`, `false` otherwise.

11.7.8 Disjointness check (a.disjoint)

Signature: `{"a.disjoint": [a, b]}`

`a`	array of any `A`
`b`	array of `A`
(returns)	boolean

Description: Return `true` if `a` and `b` are disjoint, `false` otherwise.

11.8 Functional programming

11.8.1 Map array items with function (a.map)

Signature: `{"a.map": [a, fcn]}`

`a`	array of any `A`
`fcn`	function (`A`) → any `B`
(returns)	array of `B`

Description: Apply `fcn` to each element of `a` and return an array of the results.

Details:

The order in which `fcn` is called on elements of `a` is not guaranteed, though it will be called exactly once for each element.

11.8.2 Filter array items with function (a.filter)

Signature: `{"a.filter": [a, fcn]}`

`a`	array of any `A`
`fcn`	function (`A`) → boolean
(returns)	array of `A`

Description: Apply `fcn` to each element of `a` and return an array of the elements for which `fcn` returns `true`.

Details:
The order in which `fcn` is called on elements of `a` is not guaranteed, though it will be called exactly once for each element.

11.8.3 Filter and map (a.filtermap)

Signature: `{"a.filtermap": [a, fcn]}`

`a`	array of any `A`
`fcn`	function (`A`) → union of {any `B`, null}
(returns)	array of `B`

Description: Apply `fcn` to each element of `a` and return an array of the results that are not `null`.

Details:
The order in which `fcn` is called on elements of `a` is not guaranteed, though it will be called exactly once for each element.

11.8.4 Map and flatten (a.flatmap)

Signature: `{"a.flatmap": [a, fcn]}`

`a`	array of any `A`
`fcn`	function (`A`) → array of any `B`
(returns)	array of `B`

Description: Apply `fcn` to each element of `a` and flatten the resulting arrays into a single array.

Details:
The order in which `fcn` is called on elements of `a` is not guaranteed, though it will be called exactly once for each element.

11.8.5 Reduce array items to a single value (a.reduce)

Signature: `{"a.reduce": [a, fcn]}`

`a`	array of any `A`
`fcn`	function (`A`, `A`) → `A`
(returns)	`A`

Description: Apply `fcn` to each element of `a` and accumulate a tally.

Details:

The first parameter of `fcn` is the running tally and the second parameter is an element from `a`.

The order in which `fcn` is called on elements of `a` is not guaranteed, though it accumulates from left (beginning) to right (end), called exactly once for each element. For predictable results, `fcn` should be associative. It need not be commutative.

11.8.6 Right-to-left reduce (a.reduceright)

Signature: `{"a.reduceright": [a, fcn]}`

`a`	array of any `A`
`fcn`	function (`A`, `A`) → `A`
(returns)	`A`

Description: Apply `fcn` to each element of `a` and accumulate a tally.

Details:

The first parameter of `fcn` is an element from `a` and the second parameter is the running tally.

The order in which `fcn` is called on elements of `a` is not guaranteed, though it accumulates from right (end) to left (beginning), called exactly once for each element. For predictable results, `fcn` should be associative. It need not be commutative.

11.8.7 Fold array items to another type (a.fold)

Signature: `{"a.fold": [a, zero, fcn]}`

`a`	array of any `A`
`zero`	any `B`
`fcn`	function (`B`, `A`) → `B`
(returns)	`B`

Description: Apply `fcn` to each element of `a` and accumulate a tally, starting with `zero`.

Details:

The first parameter of `fcn` is the running tally and the second parameter is an element from `a`.

The order in which `fcn` is called on elements of `a` is not guaranteed, though it accumulates from left (beginning) to right (end), called exactly once for each element. For predictable results, `fcn` should be associative with `zero` as its identity; that is, `fcn(zero, zero) = zero`. It need not be commutative.

11.8.8 Right-to-left fold (a.foldright)

Signature: `{"a.foldright": [a, zero, fcn]}`

`a`	array of any `A`
`zero`	any `B`
`fcn`	function (`B`, `A`) → `B`
(returns)	`B`

Description: Apply `fcn` to each element of `a` and accumulate a tally, starting with `zero`.

Details:

The first parameter of `fcn` is an element from `a` and the second parameter is the running tally.

The order in which `fcn` is called on elements of `a` is not guaranteed, though it accumulates from right (end) to left (beginning), called exactly once for each element. For predictable results, `fcn` should be associative with `zero` as its identity; that is, `fcn(zero, zero) = zero`. It need not be commutative.

11.8.9 Take items until predicate is false (a.takeWhile)

Signature: `{"a.takeWhile": [a, fcn]}`

`a`	array of any `A`
`fcn`	function (`A`) → boolean
(returns)	array of `A`

Description: Apply `fcn` to elements of `a` and create an array of the longest prefix that returns `true`, stopping with the first `false`.

Details:

Beyond the prefix, the number of `fcn` calls is not guaranteed.

11.8.10 Drop items until predicate is true (a.dropWhile)

Signature: `{"a.dropWhile": [a, fcn]}`

`a`	array of any `A`
`fcn`	function (`A`) → boolean
(returns)	array of `A`

Description: Apply `fcn` to elements of `a` and create an array of all elements after the longest prefix that returns `true`.

Details:

Beyond the prefix, the number of `fcn` calls is not guaranteed.

11.9 Functional tests

11.9.1 Existential check, ∃ (a.any)

Signature: `{"a.any": [a, fcn]}`

`a`	array of any `A`
`fcn`	function (`A`) → boolean
(returns)	boolean

Description: Return `true` if `fcn` is `true` for any element in `a` (logical or).

Details:

The number of `fcn` calls is not guaranteed.

11.9.2 Univeral check, ∀ (a.all)

Signature: `{"a.all": [a, fcn]}`

`a`	array of any `A`
`fcn`	function (`A`) → boolean
(returns)	boolean

Description: Return `true` if `fcn` is `true` for all elements in `a` (logical and).

Details:

The number of `fcn` calls is not guaranteed.

11.9.3 Pairwise check of two arrays (a.corresponds)

Signature: `{"a.corresponds": [a, b, fcn]}`

`a`	array of any `A`
`b`	array of any `B`
`fcn`	function (`A`, `B`) → boolean
(returns)	boolean

Description: Return `true` if `fcn` is `true` when applied to all pairs of elements, one from `a` and the other from `b` (logical relation).

Details:

The number of `fcn` calls is not guaranteed.

If the lengths of `a` and `b` are not equal, this function returns `false`.

11.10 Restructuring

11.10.1 Sliding window (a.slidingWindow)

Signature: `{"a.slidingWindow": [a, size, step, allowIncomplete]}`

`a`	array of any `A`
`size`	int
`step`	int
`allowIncomplete`	boolean
(returns)	array of array of `A`

Description: Return an array of subsequences of `a` with length `size` that slide through `a` in steps of length `step` from left to right.

Details:

If `allowIncomplete` is `true`, the last window may be smaller than `size`. If `false`, the last window may be skipped.

Runtime Errors:

If `size` is non-positive, a "size < 1" runtime error is raised.

If `step` is non-positive, a "step < 1" runtime error is raised.

11.10.2 Unique combinations of a fixed size (a.combinations)

Signature: `{"a.combinations": [a, size]}`

`a`	array of any `A`
`size`	int
(returns)	array of array of `A`

Description: Return the unique combinations of `a` with length `size`.

Runtime Errors:

If `size` is non-positive, a "size < 1" runtime error is raised.

11.10.3 Permutations (a.permutations)

Signature: `{"a.permutations": [a]}`

`a`	array of any `A`
(returns)	array of array of `A`

Description: Return the permutations of `a`.

Details:

This function scales rapidly with the length of the array. For reasonably large arrays, it will result in timeout exceptions.

11.10.4 Flatten array (a.flatten)

Signature: `{"a.flatten": [a]}`

`a`	array of array of any `A`
(returns)	array of `A`

Description: Concatenate the arrays in `a`.

11.10.5 Group items by category (a.groupby)

Signature: `{"a.groupby": [a, fcn]}`

`a`	array of any `A`
`fcn`	function (`A`) → string
(returns)	map of array of `A`

Description: Groups elements of `a` by the string that `fcn` maps them to.

12 Manipulation of other data structures

12.1 Map

12.2 Record

12.3 Enum

12.4 Fixed

13 Missing data handling

13.1 Impute library

13.1.1 Skip record (impute.errorOnNull)

Signature: `{"impute.errorOnNull": [x]}`

`x`	union of {any `A`, null}
(returns)	`A`

Description: Skip an action by raising an "encountered null" runtime error when `x` is `null`.

13.1.2 Replace with default (impute.defaultOnNull)

Signature: `{"impute.defaultOnNull": [x, default]}`

`x`	union of {any `A`, null}
`default`	`A`
(returns)	`A`

Description: Replace `null` values in `x` with `default`.

14 Aggregation

SQL-like functions

group-by tables

CUSUM

15 Descriptive statistics libraries

15.1 Sample statistics

15.1.1 Update aggregated mean (stat.sample.updateMean)

Signature: `{"stat.sample.updateMean": [runningSum, w, x]}`

`runningSum`	any record `A` with {`sum_w:` double, `sum_wx:` double}
`w`	double
`x`	double
(returns)	`A`

Description: Update a record containing running sums for computing a sample mean.

Parameters:

`runningSum`	Record of partial sums: `sum_w` is the sum of weights, `sum_wx` is the sum of weights times sample values.
`w`	Weight for this sample, which should be 1 for an unweighted mean.
`x`	Sample value.

Details:

Use `stat.sample.mean` to get the mean.

15.1.2 Compute aggregated mean (stat.sample.mean)

Signature: `{"stat.sample.mean": [runningSum]}`

`runningSum`	any record `A` with {`sum_w:` double, `sum_wx:` double}
(returns)	double

Description: Compute the mean from a `runningSum` record.

Details:

Use `stat.sample.updateMean` to fill the record.

accumulated mean, median(?)

16 Data mining models

16.1 Decision and regression Trees

16.1.1 Tree walk with simple predicates (model.tree.simpleWalk)

Signature: `{"model.tree.simpleWalk": [datum, treeNode]}`

`datum`	any record `D`
`treeNode`	any record `T` with {`field:` string, `operator:` string, `value:` any `V`, `pass:` union of {`T`, any `S`}, `fail:` union of {`T`, `S`}}
(returns)	`S`

Description: Descend through a tree comparing `datum` to each branch with a simple predicate, stopping at a leaf of type `S` (score).

Parameters:

`datum`	An element of the dataset to score with the tree.
`treeNode`	A node of the decision or regression tree. `field:` Indicates the field of `datum` to test. Fields may have any type. `operator:` One of "==" (equal), "!=" (not equal), "<" (less than), "<=" (less or equal), ">" (greater than), or ">=" (greater or equal). `value:` Value for comparison. Should be the union of or otherwise broader than all `datum` fields under consideration. `pass:` Branch to return if field `field` of `datum` `operator` `value` yields `true`. `fail:` Branch to return if field `field` of `datum` `operator` `value` yields `false`.
(return value)	The score associated with the destination leaf, which may be any type `S`. If `S` is a `string`, this is generally called a decision tree; if a `double`, it is a regression tree; if an `array` of `double`, a multivariate regression tree, etc.

Runtime Errors:

Raises a "no such field" error if `field` is not a field of `datum`.

Raises an "invalid comparison operator" error if `operator` is not one of "==", "!=", "<", "<=", ">", or ">=".

Raises a "bad value type" error if the `field` of `datum` cannot be upcast to `V`.

16.1.2 Tree walk with user-defined predicates (model.tree.predicateWalk)

Signature: `{"model.tree.predicateWalk": [datum, treeNode, predicate]}`

`datum`	any record `D`
`treeNode`	any record `T` with {`pass:` union of {`T`, any `S`}, `fail:` union of {`T`, `S`}}
`predicate`	function (`D`, `T`) → boolean
(returns)	`S`

Description: Descend through a tree comparing `datum` to each branch with a user-defined predicate,

stopping at a leaf of type `S` (score).

Parameters:

`datum`	An element of the dataset to score with the tree.
`treeNode`	A node of the decision or regression tree. `pass:` Branch to return if `"predicate": ["datum", "treeNode"]` yields `true`. `fail:` Branch to return if `"predicate": ["datum", "treeNode"]` yields `false`.
(return value)	The score associated with the destination leaf, which may be any type `S`. If `S` is a `string`, this is generally called a decision tree; if a `double`, it is a regression tree; if an `array` of `double`, a multivariate regression tree, etc.

16.2 Cluster models

16.3 Regression

16.4 Neural networks

16.5 Support vector machines